Question: Write a Producer that reads from that file, breaks the data into, and writes to Kafka and a Spark Streaming consumer that reads from the topic and emits count of rows.

Submitted by Shramana Sinha, 23f1002703

1. Implementation Steps

1.1 Google Cloud Storage Configuration

- Created a bucket named `assignmentibd`
- Uploaded an input file `input.csv` containing customer data with 1,000 records

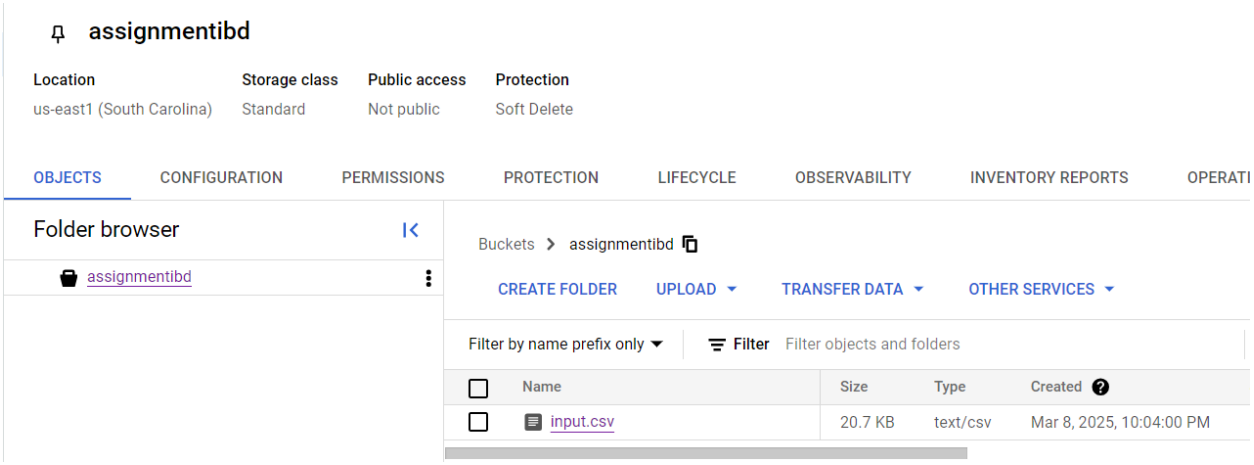

Screenshot: The gcs bucket with the input.csv file

1.2 Network Configuration

- Created a firewall rule to allow Kafka traffic on port 9092:

```
Unset
gcloud compute firewall-rules create allow-kafka --allow tcp:9092
--description="Allow Kafka traffic" --direction=INGRESS
```

1.3 Kafka Server Setup

- Launched a VM instance with Ubuntu 20.04 LTS
- Installed Java and Python dependencies:

```
Unset
sudo apt update && sudo apt install -y openjdk-8-jdk wget python3-pip
pip install kafka-python google-cloud-storage
```

- Downloaded and extracted Apache Kafka 3.9.0:

```
Unset
wget https://dlcdn.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
tar -xzf kafka_2.13-3.9.0.tgz
```

- Configured Kafka server to accept external connections:

```
Unset
# In kafka_2.13-3.9.0/config/server.properties
listeners=PLAINTEXT://0.0.0.0:9092
advertised.listeners=PLAINTEXT://<VM_EXTERNAL_IP>:9092
```

- Started Zookeeper and Kafka server as daemon processes:

```
Unset
kafka_2.13-3.9.0/bin/zookeeper-server-start.sh -daemon config/zookeeper.properties
kafka_2.13-3.9.0/bin/kafka-server-start.sh -daemon config/server.properties
```

- Uploaded the `producer.py` script and started its execution:

```
Unset
python3 producer.py
```

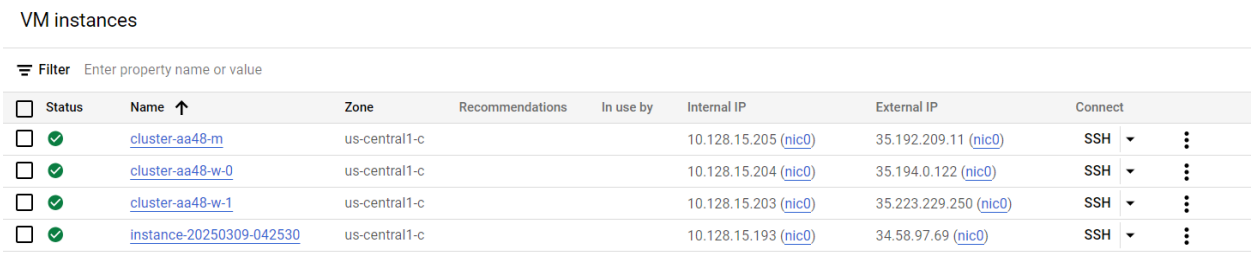

Screenshot: The vm (instance-20250309-042530) for running the kafka and producer.py

1.4 Dataproc Cluster Setup

- Created a Dataproc cluster with external IP access enabled
- Uploaded the `consumer.py` script.
- Submitted the spark job

```
Unset
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.12:3.5.4
consumer.py
```

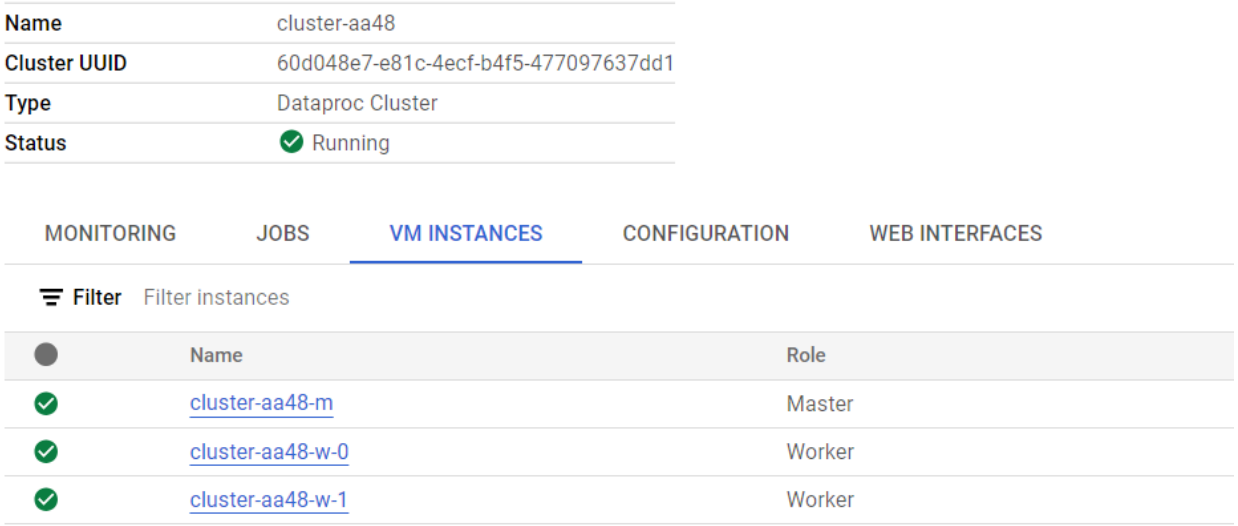

Screenshot: The dataproc cluster for running the consumer.py

2. Data Producer Implementation

The producer application performs the following tasks:

- Downloads the CSV file from Google Cloud Storage

```
Python
def download_from_gcs(bucket_name, source_blob_name, destination_file_name):
    storage_client = storage.Client()
    bucket = storage_client.bucket(bucket_name)
    blob = bucket.blob(source_blob_name)
    blob.download_to_filename(destination_file_name)
    print(
        f"Downloaded gs://{bucket_name}/{source_blob_name} to
{destination_file_name}"
    )
```

- Reads and parses the customer data

```python
def read_local_file(file_path):
    rows = []
    with open(file_path, "r") as file:
        csv_reader = csv.DictReader(file, fieldnames=["customer_id", "name",
"city"])
        for row in csv_reader:
            rows.append(row)
    return rows
```

- Streams the data to Kafka in batches of 10 records

```python
records_to_send = len(batch)
batch = batch[:records_to_send]
print(f"Sending batch {i+1} with {len(batch)} records to Kafka...")
for record in batch:
       producer.send(KAFKA_TOPIC, value=record)
       sent_count += 1
producer.flush()
```

- Implements a controlled flow with 10-second pauses between batches

```python
if i < len(batches) - 1:
       print(f"Sleeping for {sleep_seconds} seconds...")
       time.sleep(sleep_seconds)
```

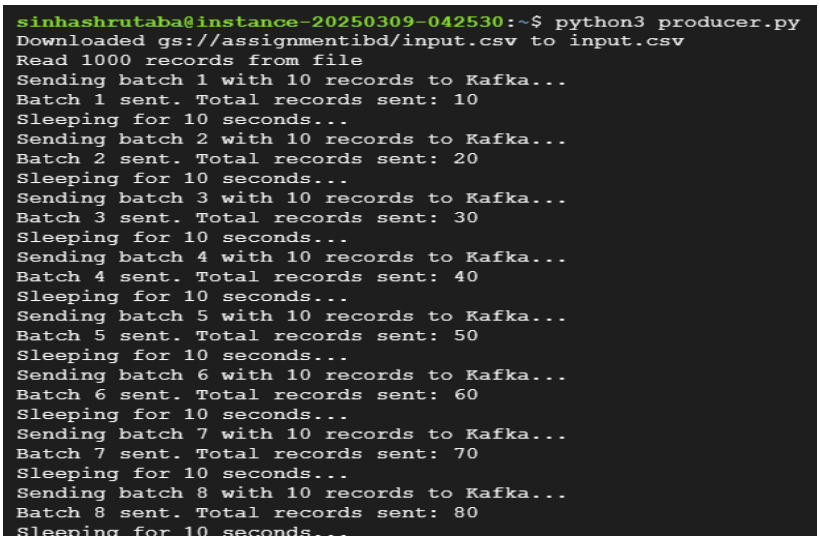

Screenshot: The producer sending batches of 10 records to kafka every 10 seconds

3. Data Consumer Implementation

The Spark Streaming consumer application performs the following functions:

- Establishes a connection to the Kafka topic and creates a streaming DataFrame from the Kafka source

```python
df = (
        spark.readStream.format("kafka")
        .option("kafka.bootstrap.servers", KAFKA_BOOTSTRAP_SERVERS)
        .option("subscribe", KAFKA_TOPIC)
        .option("startingOffsets", "earliest")
        .load()
    )
```

- Processes the incoming data by casting the binary values to string and getting the timestamp from the message.

```python
value_df = df.selectExpr("CAST(value AS STRING)", "timestamp")
```

- Counts records in 10-second windows with 5-second sliding intervals

```python
windowedCounts = value_df.groupBy(
        window(col("timestamp"), "10 seconds", "5 seconds")
    ).count()
```

- Outputs results to the console in update mode

```python
query = (
        windowedCounts.writeStream.outputMode("update")
        .format("console")
        .option("truncate", "false")
        .start()
    )
```

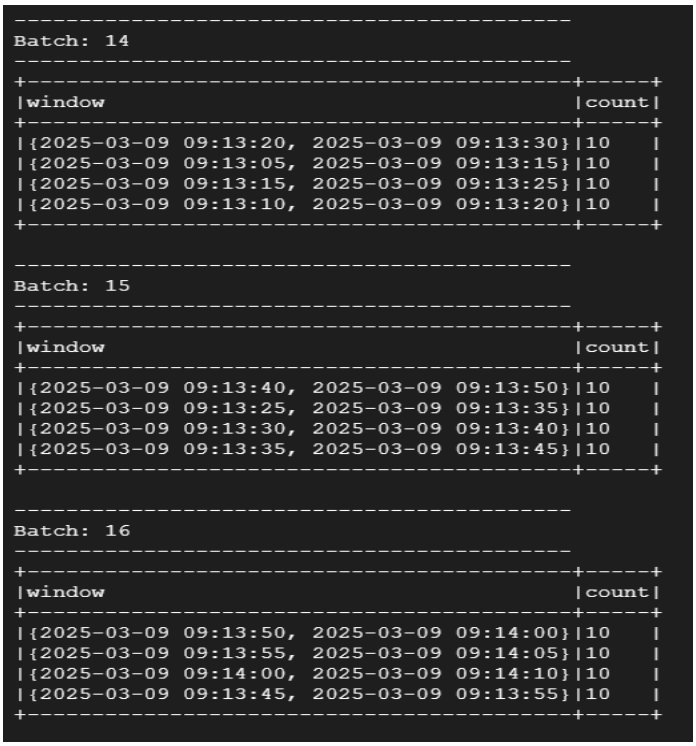

Screenshot: The output of the consumer.py showing the count of records for 10 seconds windows which are sliding every 5 seconds